

David Sammel
Final Project
CS4701

Creating Blackjack AI Using Reinforcement Learning

Blackjack

- One of the most popular casino games in the world
- Try to maximize the sum of your cards without going over 21
- Beat the dealer to win

Reinforcement Learning

- An agent observes the state of the world
- Takes an action to maximize outcome
- Tunes the decision algorithm after outcome of the action is recorded

Project Set Up

- Designed a blackjack simulation
- One dealer, and any number of players to play against the dealer
- Some players play with basic strategies, two learn as they play

Learner

- The basic learning agent is called Learner
- State Considered: Own hand, Dealer's up-facing card
- Has a matrix of probabilities of hitting given the state
- Matrix entries are initialized to .5

Learner (cont.)

- Upon winning matrix is unchanged
- Player can lose in three ways
 - Player busts: lower probability of hitting next time
 - Dealer's hand is higher: raise probability of hitting next time
 - Dealer blackjack: Don't change matrix

Counting Cards

- A way to consider more state of the game
- Keep track of the cards left to be dealt in the deck
- If a low card is dealt (2-6) increase count, if a high card is dealt (10-Ace) decrease count
- High count favors the players, low count favors the dealer
- Strategy may change depending on the count

CountingLearner

- Similar to Learner, but it considers the count as part of the state
- Has 3 matrices, one for each of when the deck is high, low, or neutral
- Plays slightly differently depending on the count

Other Players

- Dealer's playstyle
 - Another player mimics the dealer's style
 - Hits when it's hand is below 17, stands on 17 or above
- DumbPlayer
 - Randomly guesses between hitting and standing with 50% probability of both
 - Used as a baseline for the learning models

Results


- After playing 2 million hands:

Style	Dealer	Count (high)	Count (low)	Dumb	Learner	Counting Learner
Win pct.	41.4%	41.0%	42.1%	29.9%	38.3%	38.7%

Learning models perform well compared to the baseline model

Learning Rate

- Over the first 1 million hands, the model's accuracy rates:



Conclusions

- The learning agents successfully learned a rule to play blackjack
- Win percentage over time increases for both models
- Card counting model outperforms the regular model
- Winning percentage much better than the baseline, approaching dealer's strategy